

Wajahat Ahmed

AI Engineer · Machine Learning, LLMs & Computer Vision

Wajahat.ahmedhashmi@gmail.com · +923162713391 · wajahatah.github.io

Dear Zynvex Hiring Team,

I'm writing to apply for the AI Engineer role. Your description of a position sitting at the intersection of research and production maps almost exactly to how I've spent the last four years — turning research problems into deployed systems that enterprise clients depend on.

At Lambda Theta, I own AI solutions end-to-end for enterprise clients, from requirements and architecture through deployment and support. That has meant building production RAG systems and LLM agent workflows — with LangChain, LangGraph, vector databases, and LLM APIs — as FastAPI microservices backed by PostgreSQL, Redis, and Milvus, containerized with Docker and shipped through CI/CD across local GPU infrastructure and GCP. I've also fine-tuned LLMs using parameter-efficient methods (LoRA, QLoRA, PEFT) to adapt models to domain-specific tasks under real compute constraints, which speaks directly to your focus on model performance and inference latency.

On the computer-vision side, I've deployed real-time video-analytics systems processing 20+ concurrent RTSP streams with low-latency inference, and built ML pipelines spanning annotation, training, evaluation, and continuous monitoring on datasets exceeding 200K labeled samples. Optimizing models to balance accuracy, latency, and compute for production is the part of this work I enjoy most.

I've delivered all of this as part of a distributed team serving an overseas enterprise, so remote, independent, cross-functional work is simply how I operate. I'm also keen to deepen my work with Kubernetes and dedicated experiment-tracking tooling such as MLflow and Weights & Biases, building on the containerization, CI/CD, and monitoring I already run day to day.

I'd welcome the chance to discuss how I can contribute to the AI products you're building at Zynvex. Thank you for your consideration.

Best regards,

Wajahat Ahmed